

D200, Problem Set 1: Gradient Descent

Due: 4 February 2025 [here](#) in groups of 4.

Stefan Bucher

This problem set will introduce you to the gradient descent algorithm to solve the linear regression problem discussed in the lecture.¹ We will consider a univariate linear regression

$$y = \theta_0 + \theta_1 X + \varepsilon$$

and seek to fit the parameters θ_0 and θ_1 .

Problem 1

(1a) Write down the mean-squared error function and its gradient with respect to the parameters.

(1b) Implement the loss function `loss(X,y,theta)` and gradient function `gradient(X,y,theta)` in Python.

Problem 2

(2a) Implement the simple gradient descent algorithm

$$\theta_j \leftarrow \theta_j - \alpha \frac{\partial L}{\partial \theta_j}$$

as a Python function `gradient_descent(X, y, theta_init, alpha, maxsteps, precision)`. The function should take the data X , y , initial parameters θ , step size, maximum number of steps, and a precision tolerance parameter. The function should return the history of the parameters, the cost at each step, and the predictions at each step.

¹The problem set is adapted from Chi Jin's OxML 2024 Fundamentals class on Optimization.

Then, use this function to print the resulting estimate, and plot the evolution of the parameters and the mean-squared loss over 2000 iterations, with initial parameters $\theta_0 = \theta_1 = 0$ and learning rate $\alpha = 0.01$.

(2b) Compare the resulting estimates with those generated by `scipy.stats`, `statsmodels`, or `sklearn`. Plot the best-fitting line along with a scatterplot of the data.

Problem 3

Visualize the loss surface, using the provided helper functions.

```
from matplotlib import cm
import warnings
warnings.filterwarnings("ignore")

def plotting(history, cost):
    thetas1s = np.linspace(slope - 40 , slope + 30, 50)
    theta0s = np.linspace(intercept - 40 , intercept + 30, 50)
    M, B = np.meshgrid(thetas1s, theta0s)
    zs = np.array([loss(X, y, theta)
                   for theta in zip(np.ravel(B), np.ravel(M))])
    Z = zs.reshape(M.shape)
    fig = plt.figure(figsize=(10, 10))
    ax = fig.add_subplot(111, projection='3d')
    ax.plot_surface(M, B, Z, cmap=cm.coolwarm, rstride=1, cstride=1, linewidth=0, color='white')
    ax.contour(M, B, Z, 20, color='b', alpha=0.5, offset=0, stride=0)
    ax.view_init(elev=15., azim=73)
    ax.plot([history[-1][1]], [history[-1][0]], [cost[-1]] , markerfacecolor='r', markeredgewidth=2)
    ax.plot([history[0][1]], [history[0][0]], [cost[0]] , markerfacecolor='r', markeredgewidth=2)

    ax.plot([t[1] for t in history], [t[0] for t in history], cost , markerfacecolor='r', markeredgewidth=2)
    ax.plot([t[1] for t in history], [t[0] for t in history], 0 , markerfacecolor='r', markeredgewidth=2)

    ax.set_xlabel(r'$\theta_1$', fontsize=24)
    ax.set_ylabel(r'$\theta_0$', fontsize=24)
    ax.set_title(f"Iteration: {len(history)}", fontsize=24, fontweight='bold')

def plotting_diverge(history, cost):
    #print("X: ", theta[1], " abs(theta[1]) :", abs(theta[1]))
    thetas1s = np.linspace( -(-0.34 + abs(theta[1])*0.6 + 50) , -0.34 + abs(theta[1])*0.6 + 50, 50)
```

```

theta0s = np.linspace(79.18 - abs(theta[1]) -200, 79.18 + abs(theta[1]) +200, 50)
M, B = np.meshgrid(theta1s, theta0s)
zs = np.array([loss(X, y, theta)
                for theta in zip(np.ravel(B), np.ravel(M))])
Z = zs.reshape(M.shape)
fig = plt.figure(figsize=(10, 10))
ax = fig.add_subplot(111, projection='3d')
ax.plot_surface(M, B, Z, cmap=cm.coolwarm, rstride=1, cstride=1, color='b', alpha=0.8)
ax.contour(M, B, Z, 20, color='b', alpha=0.5, offset=0, stride=0)
ax.view_init(elev=15., azimuth=73)
ax.plot([history[0][1]], [history[0][0]], [cost[0]], markerfacecolor='r', markeredgecolor='b')
ax.plot([history[0][1]], [history[0][0]], [cost[0]], markerfacecolor='r', markeredgecolor='b')

ax.plot([t[1] for t in history], [t[0] for t in history], cost, markerfacecolor='r', markeredgecolor='b')
ax.plot([t[1] for t in history], [t[0] for t in history], 0, markerfacecolor='r', markeredgecolor='b')

ax.set_xlabel(r'$\theta_1$', fontsize=24)
ax.set_ylabel(r'$\theta_0$', fontsize=24)
ax.set_title(f"Iteration: {len(history)}", fontsize=24, fontweight='bold')

def gen_gif(method_name,h,c,div=False):
    %matplotlib notebook
    from tqdm import tqdm
    #import os
    # generate gif, run after the specific method run
    theta = h[-1]
    method_name = method_name

    # use tqdm to create a progress bar
    print("Produce the pngs")
    for i in range(1,len(h)+1):
        history_, cost_ = h[0:i], c[0:i]
        if not div:
            plotting(history_,cost_)
        else:
            plotting_diverge(history_,cost_)
        plt.savefig(f'figures/{method_name}-{i}.png')

import imageio
import os
print("Produce the GIF")

```

```

with imageio.get_writer(f'figures/{method_name}.gif', mode='I') as writer:
    for filename in tqdm([f'figures/{method_name}-{i}.png' for i in range(1, len(h)+1)], desc='writing'):
        image = imageio.imread(filename)
        writer.append_data(image)
        os.remove(filename)

def plotting_loss_surface():
    thetas = np.linspace(slope - 40, slope + 30, 50)
    theta0s = np.linspace(intercept - 40, intercept + 30, 50)
    M, B = np.meshgrid(thetas, theta0s)
    zs = np.array([loss(X, y, theta)
                   for theta in zip(np.ravel(B), np.ravel(M))])
    Z = zs.reshape(M.shape)
    fig = plt.figure(figsize=(10, 10))
    ax = fig.add_subplot(111, projection='3d')
    ax.plot_surface(M, B, Z, cmap=cm.coolwarm, rstride=1, cstride=1, linewidth=0, color='b',
                   alpha=0.5, offset=0, stride=0)
    ax.view_init(elev=15., azim=73)

    ax.set_xlabel(r'$\theta_1$', fontsize=24)
    ax.set_ylabel(r'$\theta_0$', fontsize=24)

```

Use function `plotting(history, costs)` to visualize the loss surface for various values of the hyperparameter `alpha` to get an intuition for how the learning rate affects the performance of the algorithm. Discuss your observations.